Interface Control Document (ICD)
Data Readiness System (DRS) / Data Readiness System (ADRS)

ICD v1.0 (MVP Baseline)
Organization: The AI Institute Africa
Date: 19 February 2026

Document Control

Document Name	ICD v1.0 - Data Readiness System (DRS)
Version	1.0
Status	Draft for MVP implementation
Owner	Project Lead (Wills) / Delivery Team
Primary Users	Backend developers, frontend developers, QA/testers, integration partners
Precedence	CDM v1 Schema takes precedence for canonical fields; MVP Scope defines in-scope interfaces

1. Purpose and Scope

This Interface Control Document (ICD) defines the API and data contract baseline for the DRS MVP. It is intended to be used by the delivery team using AI Coding platforms to implement, test, and verify integrations across DRS subsystems, including optional integration with an external ingestion platform.

The MVP must support: evidence ingestion and immutable storage, OCR/extraction, trust scoring and human validation, mapping approved values into the CDM, basic entity resolution into golden records, dataset publishing with dataset cards, and governance via RBAC and audit logs.

1.1 In-Scope Interfaces (MVP)

- DRS REST API (Gateway) for UI and integrations
- Evidence ingestion (native upload + external platform registration mode)
- OCR/extraction job triggering and status
- Validation task queue and approve/reject actions
- CDM mapping execution and canonical record retrieval
- Basic entity resolution and golden record retrieval

- Dataset publishing, dataset version retrieval, download, and dataset cards
- Audit log retrieval for traceability and governance

1.2 Out-of-Scope (MVP)

- Production-grade connectors for SharePoint/Google Drive (beyond optional external platform REST integration)
- Advanced streaming ingestion and large-scale distributed optimization
- Automated de-identification/pseudonymization engine (unless explicitly added later)
- Advanced knowledge-graph reasoning engines

2. Architectural Context (for Interfaces)

The DRS MVP follows a layered pipeline architecture: Evidence Layer → Intelligence Layer → Canonical Data Layer → Knowledge Layer → Publishing Layer, with Governance controls applied across all layers.

2.1 Primary Actor Roles (RBAC)

The MVP supports the following roles and minimum responsibilities:

Role	Primary Responsibilities	Typical API Capabilities
Archive Operator	Create batches (optional), upload/register evidence	Upload evidence, view own evidence
Data Validator	Validate extracted fields	List/claim tasks, approve/reject
AI Engineer	Configure extraction/mapping; run mapping	Trigger extraction, run mapping, view canonical
Compliance Officer	Approve publishing; review audit logs	Approve publish, view audit, view datasets
Admin	Manage users/roles/tenants/config	User/role management, configuration

3. API Conventions

3.1 Base URL and Versioning

Base path: /api/v1 (all endpoints in this ICD are scoped under this prefix).

Schema versioning: CDM schema_version is required in publishing and dataset cards (semantic versioning).

3.2 Authentication and Authorization

Authentication: OAuth2/OpenID Connect or JWT (MVP minimum: JWT).

Authorization: RBAC enforced on all endpoints; each endpoint specifies minimum role(s).

3.3 Common Headers

Header	Usage
Authorization	Bearer <token>
Content-Type	application/json (unless multipart upload or file download)
X-Tenant-Id	Required if the token is not tenant-scoped; otherwise optional
Idempotency-Key	Required for POST /evidence/upload and POST /evidence/register-external
X-Request-Id	Optional client-provided correlation id; echoed in responses/logs

3.4 Standard Response Envelope

Unless otherwise specified, JSON responses follow:

```
{
  "success": true,
  "data": <object|array|null>,
  "error": null,
  "request_id": "uuid"
}
```

3.5 Standard Error Object

```
{
  "code": "VALIDATION_ERROR|NOT_FOUND|FORBIDDEN|CONFLICT|INTERNAL",
  "message": "Human readable summary",
  "details": [{"field":"...", "issue":"..."}]
}
```

3.6 Pagination

List endpoints support: ?limit=50&cursor=<opaque>. Response includes next_cursor when available.

3.7 Timestamps and IDs

Timestamps are ISO-8601 UTC (e.g., 2026-02-16T12:00:00Z).

IDs are opaque strings (e.g., EVID-98A2) or UUIDs; clients must treat them as opaque.

3.8 Evidence Pointer Format (CDM precedence)

Evidence pointers use page number and bounding box coordinates. Pointer schema is defined in Section 4. All extracted and canonical values that are evidence-backed should reference an EvidencePointer.

4. Data Contracts (Core Schemas)

This section defines the MVP baseline payload schemas. Canonical entity fields follow CDM v1 and must take precedence in implementation and testing.

4.1 EvidenceObject

Field	Type	Required	Notes
evidence_id	string	Yes	DRS identifier
tenant_id	string	Yes	Tenant namespace
batch_id	string	No	Optional grouping (MVP supports batch/folder)
folder_id	string	No	Optional grouping
source_type	string	Yes	SCAN|UPLOAD|EXTERNAL
source_reference	string	No	e.g., batch/folder path or external reference
file_name	string	No	Original filename
file_hash	string	Yes	sha256:<hex>
file_format	string	Yes	pdf|png|jpg|tiff|...
file_uri	string	Yes	Object store URI or external URI
ingested_at	string(datetime)	Yes	UTC
ingested_by	string	Yes	User id
document_type_hint	string	No	Optional hint at ingestion
status	string	Yes	INGESTED|PROCESSING|PROCESSED|ARCHIVED
immutability_status	string	Yes	LOCKED once stored

4.2 Document (CDM)

Field	Type	Required	Notes
document_id	string	Yes	PK
tenant_id	string	Yes	
evidence_id	string	Yes	FK → EvidenceObject
title	string	No	
document_date	string(date)	No	
source_system	string	No	or source_type
source_type	string	No	SCAN|UPLOAD|EXTERNAL
ingested_at	string(datetime)	Yes	

ingested_by	string	Yes	
document_type_hint	string	No	From ingestion
classification_label	string	No	e.g., Invoice
classification_confidence	number	No	0-1
status	string	Yes	INGESTED|PROCESSED|VALIDATED

4.3 EvidencePointer (CDM)

Field	Type	Required	Notes
pointer_id	string	Yes	PK
evidence_id	string	Yes	FK → EvidenceObject
page_number	integer	Yes	0-based or 1-based (choose one; MVP recommends 1-based)
bbox_x1	number	Yes	Top-left X
bbox_y1	number	Yes	Top-left Y
bbox_x2	number	Yes	Bottom-right X
bbox_y2	number	Yes	Bottom-right Y
snippet_text	string	No	Optional snippet
pointer_format_version	string	Yes	e.g., v1

4.4 Attribute (CDM)

Field	Type	Required	Notes
attribute_id	string	Yes	PK
tenant_id	string	Yes	
subject_ref_type	string	Yes	OBJECT|PARTY|EVENT|DOCUMENT
subject_ref_id	string	Yes	
attribute_key	string	Yes	Canonical key (snake_case)
value_type	string	Yes	string|number|date|boolean|json
value_raw	string	Yes	Raw extracted value
value_normalized	string	No	Optional normalized value
unit	string	No	
currency	string	No	
confidence_score	number	Yes	0-1
evidence_pointer_id	string	Yes	FK → EvidencePointer
validation_state	string	Yes	PENDING|APPROVED|REJECTED
validated_by	string	No	
validated_at	string(datetime)	No	

4.5 Party (CDM)

Field	Type	Required	Notes
party_id	string	Yes	PK
tenant_id	string	Yes	

party_kind	string	Yes	PERSON|ORG
display_name	string	Yes	
primary_identifier_id	string	No	FK → Identifier
golden_record_id	string	No	FK → GoldenRecord
created_at	string(datetime)	Yes	
updated_at	string(datetime)	Yes	

4.6 Identifier (CDM)

Field	Type	Required	Notes
identifier_id	string	Yes	PK
tenant_id	string	Yes	
subject_ref_type	string	Yes	PARTY|OBJECT|EVENT
subject_ref_id	string	Yes	
id_type_label	string	Yes	taxonomy label (NationalID, Passport, CompanyRegNo, etc.)
id_value	string	Yes	
issuing_authority	string	No	
is_verified	boolean	Yes	
evidence_pointer_id	string	No	FK → EvidencePointer

4.7 Event (CDM)

Field	Type	Required	Notes
event_id	string	Yes	PK
tenant_id	string	Yes	
event_time	string(datetime|date)	No	
event_type_label	string	Yes	taxonomy label
event_confidence	number	No	0-1
source_document_id	string	Yes	FK → Document
status	string	Yes	CANDIDATE|VALIDATED|PUBLISHED
created_at	string(datetime)	Yes	
updated_at	string(datetime)	Yes	

4.8 Object (CDM)

Field	Type	Required	Notes
object_id	string	Yes	PK
tenant_id	string	Yes	
object_type_label	string	Yes	taxonomy label (Invoice, Policy, Case, etc.)
object_type_confidence	number	No	0-1
source_document_id	string	Yes	FK → Document
status	string	Yes	CANDIDATE|VALIDATED|RESOLVED|PUBLIS

			HED
created_at	string(datetime)	Yes	
updated_at	string(datetime)	Yes	

4.9 Relationship (CDM)

Field	Type	Required	Notes
relationship_id	string	Yes	PK
tenant_id	string	Yes	
from_ref_type	string	Yes	PARTY|EVENT|DOCUMENT|OBJECT
from_ref_id	string	Yes	
to_ref_type	string	Yes	
to_ref_id	string	Yes	
relationship_type_label	string	Yes	taxonomy label (PAID_TO, OWNS, etc.)
confidence	number	No	0-1
source_document_id	string	No	FK → Document
created_at	string(datetime)	Yes	

4.10 GoldenRecord (CDM)

Field	Type	Required	Notes
golden_record_id	string	Yes	PK
tenant_id	string	Yes	
subject_kind	string	Yes	PARTY|OBJECT
canonical_subject_id	string	Yes	Surviving subject id
merged_from_ids	array[string]	Yes	List of merged ids
merge_rule_version	string	Yes	
match_score	number	No	0-1
merge_decision_log_ref	string	No	Audit reference

4.11 Standard Record Object (SRO)

Field	Type	Required	Notes
sro_id	string	Yes	PK
tenant_id	string	Yes	
evidence_id	string	Yes	
document_id	string	No	FK → Document when created
extraction_run_id	string	Yes	
candidate_fields	array[CandidateField]	Yes	See CandidateField
trust_score	number	No	0-1

status	string	Yes	PENDING_VALIDATION|APPROVED|REJECTED
created_at	string(datetime)	Yes	

4.12 CandidateField

Field	Type	Required	Notes
field_key	string	Yes	Canonical key (snake_case)
value_raw	string	Yes	
value_type	string	Yes	string|number|date|boolean|json
confidence_score	number	Yes	0-1
evidence_pointer_id	string	Yes	FK → EvidencePointer
pii_flag	boolean	No	Optional governance signal
validation_state	string	Yes	PENDING|APPROVED|REJECTED
validated_value	string	No	Set on approval

4.13 ValidationTask

Field	Type	Required	Notes
validation_id	string	Yes	PK
tenant_id	string	Yes	
sro_id	string	Yes	FK → SRO
assigned_to	string	No	Validator user id
status	string	Yes	PENDING_VALIDATION|APPROVED|REJECTED|NEEDS_RESCAN|ESCALATED
fields_to_validate	array[string]	Yes	List of field_key values
created_at	string(datetime)	Yes	
updated_at	string(datetime)	Yes	

4.14 DatasetVersionMetadata (CDM)

Field	Type	Required	Notes
dataset_version_id	string	Yes	PK
tenant_id	string	Yes	
dataset_name	string	Yes	
dataset_version	string	Yes	semver e.g., v1.0.0
schema_version	string	Yes	CDM schema version
generated_at	string(datetime)	Yes	
lineage_evidence_batches	array[string]	Yes	Source batch ids
validation_summary	string	Yes	Counts, approvals
approvals	array[Approval]	No	Optional for MVP
download_uris	object	No	csv/json/parquet

			URIs

4.15 AuditEvent

Field	Type	Required	Notes
event_id	string	Yes	PK
tenant_id	string	Yes	
actor	string	Yes	User id
action	string	Yes	See Appendix A
object_ref	string	Yes	e.g., EVID-... or SRO-...
timestamp	string(datetime)	Yes	
before	object	No	Optional
after	object	No	Optional
reason	string	No	Optional

5. REST API Endpoints (Gateway)

All endpoints are prefixed with /api/v1. Minimum roles are listed for each endpoint.

5.1 Evidence Ingestion

Upload Evidence (Native Ingestion)

Method	POST
Path	/evidence/upload
Min Roles	Archive Operator, Admin
Description	Uploads a file and creates an immutable EvidenceObject. Returns evidence_id and stored file_uri. Must be idempotent.
Notes	On duplicate Idempotency-Key, return the original success response. Compute SHA-256 and lock evidence after write.

Request (example):

Content-Type: multipart/form-data
Fields:
- file: <binary>
- tenant_id (optional if token-scoped)
- batch_id (optional)
- folder_id (optional)
- document_type_hint (optional)
- source_reference (optional)

Headers:
- Idempotency-Key: <uuid>

Response (example):

{ "success": true, "data": { "evidence_id":"EVID-98A2", "file_hash":"sha256:...", "file_uri":"s3://evidence/.../original.pdf", "immutability_status":"LOCKED" }, "error": null }

Register Evidence (External Ingestion Platform Mode)

Method	POST
Path	/evidence/register-external
Min Roles	Archive Operator, Admin
Description	Registers an evidence file stored in an external ingestion platform (Mode B). DRS stores provenance and external references.
Notes	Requires Idempotency-Key. DRS may fetch metadata/hash via external platform API as part of registration.

Request (example):

{
  "external_platform_id": "EXT-PLATFORM-001",
  "external_file_id": "FILE-12345",
  "tenant_id": "TENANT-001",
  "batch_id": "BATCH-2026-0001",
  "document_type_hint": "Invoice"
}

Response (example):

{ "success": true, "data": { "evidence_id":"EVID-EXT-001", "source_type":"EXTERNAL", "external_file_id":"FILE-12345", "status":"INGESTED" }, "error": null }

Get Evidence Metadata

Method	GET
Path	/evidence/{evidence_id}/metadata
Min Roles	Archive Operator, Data Validator, AI Engineer, Compliance Officer, Admin (RBAC enforced)
Description	Retrieves EvidenceObject metadata (no file

	content).
Notes	Access must be logged (AUDIT: ACCESS_EVIDENCE_METADATA).

Response (example):

{ "success": true, "data": { "evidence_id":"EVID-98A2", "file_hash":"sha256:...", "file_format":"pdf", "ingested_at":"2026-02-16T12:00:00Z", "status":"INGESTED" }, "error": null }

Download Evidence File

Method	GET
Path	/evidence/{evidence_id}/download
Min Roles	Archive Operator, Data Validator, Compliance Officer, Admin (RBAC enforced)
Description	Downloads the evidence file or returns a signed URL.
Notes	Access must be logged (AUDIT: DOWNLOAD_EVIDENCE).

Response (example):

Binary stream (application/pdf) OR { signed_url: "...", expires_at: "..." }

5.2 Document Intelligence (OCR/Extraction)

Trigger OCR + Extraction

Method	POST
Path	/intelligence/extract
Min Roles	AI Engineer, Admin
Description	Creates an extraction run for an evidence_id (and optional document_id). Produces SRO + evidence pointers.
Notes	Must write audit events OCR_RUN and EXTRACTION_RUN when jobs actually execute.

Request (example):

{
  "evidence_id": "EVID-98A2",
  "document_type_hint": "Invoice",
  "options": { "run_ocr": true, "run_tables": false }
}

Response (example):

{ "success": true, "data": { "extraction_run_id":"EXTR-001", "sro_id":"SRO-919", "status":"QUEUED" }, "error": null }

Get Extraction Run Status

Method	GET
Path	/intelligence/extract/{extraction_run_id}
Min Roles	AI Engineer, Admin
Description	Returns current status and summary (pages processed, errors).
Notes	Status: QUEUED|RUNNING|COMPLETED|FAILED

Response (example):

{ "success": true, "data": { "extraction_run_id":"EXTR-001", "status":"COMPLETED", "created_sro_id":"SRO-919" }, "error": null }

Get SRO (Candidate Extraction Output)

Method	GET
Path	/records/{sro_id}
Min Roles	Data Validator, AI Engineer, Admin
Description	Retrieves the Standard Record Object with candidate fields, confidence scores, and evidence pointers.
Notes	Used by Validation UI.

Response (example):

{ "success": true, "data": { "sro_id":"SRO-919", "candidate_fields":[{"field_key":"invoice_number","value_raw":"3312","confidence_score":0.91,"evidence_pointer_id":"PTR-1","validation_state":"PENDING"}] }, "error": null }

5.3 Trust and Validation (HITL)

List Validation Tasks

Method	GET
Path	/validation/tasks?status=PENDING_VALIDATION&limit=50
Min Roles	Data Validator, Admin
Description	Returns validation tasks visible to the caller. Supports filtering by status, assigned_to, batch_id.
Notes	When tasks are created, write audit event CREATE_VALIDATION_TASK.

Response (example):

{ "success": true, "data": { "items":[{"validation_id":"VAL-002","sro_id":"SRO-919","status":"PENDING_VALIDATION"}],"next_cursor": null }, "error": null }

Assign (Claim) Validation Task

Method	POST
Path	/validation/tasks/{validation_id}/claim
Min Roles	Data Validator, Admin
Description	Claims a task for the current validator (sets assigned_to).
Notes	Conflict if already assigned to another user (HTTP 409).

Response (example):

{ "success": true, "data": { "validation_id":"VAL-002", "assigned_to":"validator_007" }, "error": null }

Approve Field(s)

Method	POST
Path	/validation/approve
Min Roles	Data Validator, Admin
Description	Approves one or more candidate fields on an SRO; stores validated_value and audit trail.
Notes	Write audit event APPROVE_FIELD per field. If all required fields are approved, task may transition to APPROVED.

Request (example):


```
{
  "sro_id": "SRO-919",
  "approvals": [
    {"field_key":"invoice_number","validated_value":"3312"},
    {"field_key":"total_amount","validated_value":"1000.00"}
  ]
}
```


Response (example):


```
{ "success": true, "data": { "sro_id":"SRO-919", "status":"PENDING_VALIDATION" }, "error": null }
```


Reject Field(s)

Method	POST
Path	/validation/reject
Min Roles	Data Validator, Admin
Description	Rejects one or more candidate fields; optional reason. May trigger rescan/escalation.
Notes	Write audit event REJECT_FIELD per field.

Request (example):


```
{
  "sro_id": "SRO-919",
  "rejections": [
    {"field_key":"supplier_name","reason":"Illegible"}
  ],
  "next_action": "NEEDS_RESCAN"
}
```


Response (example):


```
{ "success": true, "data": { "sro_id":"SRO-919", "status":"NEEDS_RESCAN" }, "error": null }
```


5.4 Canonical Data (CDM) Mapping

Map Approved SRO into CDM

Method	POST

Path	/cdm/map
Min Roles	AI Engineer, Admin
Description	Maps validated fields from an SRO into CDM entities (Document, Party, Event, Object, Attribute, Relationship, Identifier).
Notes	Write audit event MAP_TO_CDM. Reject mapping if required validations are incomplete.

Request (example):

{ "sro_id":"SRO-919", "schema_version":"v1.0.0" }

Response (example):

{ "success": true, "data": { "mapping_run_id":"MAP-001", "created": { "document_id":"DOC-111", "party_ids":["PTY-1"], "object_ids":["OBJ-1"], "attribute_count": 12 } }, "error": null }

Get Canonical Document

Method	GET
Path	/cdm/documents/{document_id}
Min Roles	AI Engineer, Data Validator, Compliance Officer, Admin
Description	Returns CDM Document record with links to evidence and related objects.
Notes	Used by Canonical Record View screen.

Response (example):

{ "success": true, "data": { "document_id":"DOC-111", "evidence_id":"EVID-98A2" }, "error": null }

List Attributes for a Subject

Method	GET
Path	/cdm/attributes?subject_ref_type=OBJECT&subject_ref_id=OBJ-1
Min Roles	AI Engineer, Compliance Officer, Admin
Description	Lists evidence-backed Attributes for a specific subject (Object/Party/Event/Document).
Notes	Supports pagination.

Response (example):

{ "success": true, "data": { "items":[{"attribute_key":"invoice_number","value_raw":"3312","validation_state":"APPROVED"}] }, "error": null }

5.5 Entity Resolution (Basic MVP)

Run Entity Resolution (Parties)

Method	POST
Path	/resolution/parties/run
Min Roles	AI Engineer, Admin
Description	Runs basic matching rules for Party entities and creates/updates GoldenRecords.
Notes	Write audit event MERGE_ENTITIES when merges occur.

Request (example):

{ "party_ids":["PTY-1","PTY-2"], "merge_rule_version":"v1" }

Response (example):

{ "success": true, "data": { "run_id":"RES-001", "golden_record_ids":["GR-1"] }, "error": null }

Get Golden Record

Method	GET
Path	/resolution/golden-records/{golden_record_id}
Min Roles	AI Engineer, Compliance Officer, Admin
Description	Retrieves a GoldenRecord and merged_from_ids.
Notes	MVP may store golden records in relational DB (graph optional).

Response (example):

{ "success": true, "data": { "golden_record_id":"GR-1", "merged_from_ids":["PTY-1","PTY-2"] }, "error": null }

5.6 Dataset Publishing

Publish Dataset

Method	POST
Path	/publish/dataset
Min Roles	Compliance Officer, Admin
Description	Generates a dataset export from canonical records, creates a DatasetVersionMetadata record, and stores dataset card.
Notes	Write audit event PUBLISH_DATASET. Enforce approval rules if configured.

Request (example):


```
{
  "dataset_name": "finance_dataset",
  "schema_version": "v1.0.0",
  "source_batches": ["BATCH-2026-0001"],
  "formats": ["CSV","JSON"]
}
```


Response (example):

{ "success": true, "data": { "dataset_version_id":"DSV-001", "dataset_version":"v1.0.0", "generated_at":"2026-02-16T13:00:00Z" }, "error": null }

List Dataset Versions

Method	GET
Path	/publish/datasets?limit=50
Min Roles	Compliance Officer, AI Engineer, Admin
Description	Lists published dataset versions for the tenant.
Notes	Supports pagination.

Response (example):

{ "success": true, "data": { "items":[{"dataset_version_id":"DSV-001","dataset_name":"finance_dataset","dataset_version":"v1.0.0"}] }, "error": null }

Download Dataset File

Method	GET
Path	/publish/dataset/{dataset_version_id}/download?format=CSV
Min Roles	Compliance Officer, AI Engineer, Admin
Description	Downloads the dataset export file (or returns signed URL).
Notes	Write audit event DOWNLOAD_DATASET.

Response (example):

Binary stream OR { signed_url: "..." }

Get Dataset Card

Method	GET
Path	/publish/dataset/{dataset_version_id}/card
Min Roles	Compliance Officer, AI Engineer, Admin
Description	Retrieves dataset card (minimum fields per MVP scope).
Notes	Dataset card must include lineage and validation summary.

Response (example):

{ "success": true, "data": { "dataset_name":"finance_dataset","schema_version":"v1.0.0","lineage_evidence_batches":["BATCH-2026-0001"] }, "error": null }

5.7 Governance and Audit

List Audit Events

Method	GET
Path	/audit/events?limit=100&cursor=<cursor>
Min Roles	Compliance Officer, Admin
Description	Lists audit events (filterable by action, actor, object_ref, time range).
Notes	Mandatory audit events are listed in Appendix A.

Response (example):

{ "success": true, "data": { "items":[{"event_id":"AUD-778","action":"APPROVE_FIELD","object_ref":"SRO-919.total_amount"}], "next_cursor": null }, "error": null }

5.8 External Ingestion Platform (Optional Integration)

The MVP supports integrating with an external system that provides file ingestion/storage/tracking. DRS may call external APIs for metadata, download, and status; DRS also exposes internal endpoints to support registration and processing.

Sync External File Status

Method	POST
Path	/integrations/external/{external_platform_id}/files/{external_file_id}/sync
Min Roles	AI Engineer, Admin
Description	Fetches latest status/metadata from the external platform and updates the linked EvidenceObject.
Notes	Relies on external platform contract defined in Section 6.

Response (example):

{ "success": true, "data": { "external_tracking_status":"ACTIVE", "sync_timestamp":"2026-02-16T12:10:00Z" }, "error": null }

6. External Ingestion Platform API Contract (Mode B)

If Mode B is enabled, the external platform must expose the following minimum endpoints. These are the contract assumptions for integration testing.

External API	Method/Path	Minimum Response Fields	Notes
File Metadata Fetch	GET /external/files/{file_id}/metadata	file_name, created_at, file_type, storage_uri or download_link, file_hash (optional)	DRS uses for evidence registration and verification
File Download Access	GET /external/files/{file_id}/download	file stream or signed URL	Used when DRS needs to run OCR/extraction
File Tracking Status	GET /external/files/{file_id}/status	status: ACTIVE|ARCHIVED|DELETED	DRS stores external_tracking_status and sync_timestamp

6.1 DRS Stored Fields for External Integration

When integrated, DRS stores the following on EvidenceObject (or linked integration table):

- external_platform_id
- external_file_id
- external_tracking_status
- sync_timestamp
- external_uri

7. Audit Event Taxonomy (MVP Mandatory)

The following audit actions are mandatory for MVP (logged in tamper-evident storage):

- UPLOAD_EVIDENCE
- OCR_RUN
- EXTRACTION_RUN
- CREATE_VALIDATION_TASK
- APPROVE_FIELD
- REJECT_FIELD
- MAP_TO_CDM
- MERGE_ENTITIES
- PUBLISH_DATASET
- DOWNLOAD_DATASET

8. Integration Test Flows (for MVP)

8.1 Native Evidence Flow (Mode A)

1. POST /evidence/upload → evidence_id
2. POST /intelligence/extract (evidence_id) → sro_id
3. GET /validation/tasks → validation_id
4. POST /validation/approve and/or /validation/reject until task complete
5. POST /cdm/map (sro_id) → canonical ids
6. POST /resolution/parties/run (optional basic MVP) → golden_record_id
7. POST /publish/dataset → dataset_version_id
8. GET /publish/dataset/{dataset_version_id}/card and /download
9. GET /audit/events to verify traceability

8.2 External Evidence Flow (Mode B)

10. POST /evidence/register-external (external_file_id) → evidence_id
11. POST /integrations/external/{platform}/files/{file}/sync (optional)
12. POST /intelligence/extract (evidence_id) → sro_id (DRS downloads from external platform)

13. Continue as in Mode A from validation onward

8.3 Traceability Assertions (MVP)

- Every approved Attribute must have evidence_pointer_id linking back to EvidencePointer → evidence_id
- Dataset card must list lineage_evidence_batches and schema_version
- Audit events must show: who did what, when, and on which object (evidence, SRO field, dataset version)

9. References

- Detailed Design Document (DDD) - Data Readiness System (DRS) v1.0, February 2026
- System Requirements Specification (SRS) - Data Readiness System (DRS) v1.0
- MVP Scope Document (MVP-SCOPE v1.1), February 2026
- CDM v1 Schema (CDM v1.x) - Common Data Model baseline for MVP